# Data Compression Algorithm Design

**Author:** *Isaac Shiells Thomas*

**Email:** isaacshiellsthomas.work@proton.me

## Contents

## 1 Context

This is an algorithm designed to compress a given buffer of data in the form of a byte array.

The algorithm lives within a function. This function is called with two arguments; a pointer to a data buffer in the form of a byte array (`data_ptr`), and the number of bytes to compress in the form of an integer (`data_size`). After the function executes, the data in the buffer will be modified and the size of the modified buffer will be returned.

It is not expected for the algorithm to reallocate memory for the new size of the buffer. The space after the requested number of bytes to be compressed will not be altered, and the trailing space — the space saved by the compression — will be zeroed.

### 1.1 Assumptions

1. The `data_ptr` will point to an array of bytes. Each byte will contain a number from 0 to 127 (`0x00` to `0x7F`). Additionally, it is common for the data in the buffer to have the same value repeated multiple times in a row.
2. The compressed data needs to be decompressable. A accompanying function that will decompress the data must exist.

### 1.2 Goals

The first goal of this algorithm will be to compress the data by utilizing the fact that it can be commonly expected that bytes will appear multiple times sequentially. The main tactic will be to condense repeated sequential bytes into an encoding of two bytes where the first byte is the original value and the second byte is the number of times in a row that byte was repeated.

The second goal will be that, even in the worse case scenario, this algorithm will **not** increase the length of the data and at worst will result in a compressed data length equal to the inputted data length (`data_size`) when measured in bytes.

## 2 Terms and Definitions

**Sequential Byte**   A *sequential byte* is a set of bytes in a stream that appears twice or more in a row with the exact same value. An instance of a byte that appears twice or more in a row is called *sequential*.

**Byte Multiplicity**  When an instance of a byte is sequential, its *multiplicity* is the *number of times* that byte appears in a row. In the byte stream (`0x64`, `0x4A`, `0x64`, `0x64`, `0x64`), the second instance of the byte `0x64` has a *multiplicity* of **3**. A byte that is *not sequential* has a *multiplicity* of **1**.

**Semi-Unique Byte Set**  A set of bytes where order matters and each byte may appear more than once, but not more than once in a row.

# 3 Algorithm

This algorithm will compare chunks of 7 *semi-unique* bytes at a time, starting from the $0^{\text{th}}$ byte of the byte array. First, the initial byte is recorded, the algorithm then scans each subsequent byte until a byte of a different value is found. At that point, the quantity of the first byte, and it's value, are recorded as the *first* of the *7-byte-chunk*, and then the algorithm continues for the next byte (the one that ended the sequential streak of the first byte). The process is repeated until 7 bytes are found. If the algorithm comes across a byte that is already present as a previous value in the *7-byte-chunk*, then it treats it *no differently* than if it was completely new. The algorithm only cares about separating and condensing the sequential streams of bytes (bytes that all have the same value and appear in a row). It does not care if the *7-byte-chunk* consists of only two different bytes, alternating back and forth.

Now that the *7-byte-chunk* has been filled, the algorithm will handle two cases differently: 1) When their exists at least one byte that appears three or more times in a row, and 2) when there doesn't.

In both cases, a header byte will be used to facilitate the decoding process. Case 1 will use a full header byte while case 2 will use a partial header byte (details explained below). Despite their differences, the two cases use a single high order bit to declare to the decoder which type of compression was used for the following block of bytes. A high order bit of value "0" indicates that the sub-algorithm defined in case 1 is used, while a high order bit of value "1" indicates that the sub-algorithm defined in case 2 is used.

The two cased are defined as follows:

## 3.1 Case 1: 3 or More Sequential Bytes Found

This case occurs when at least one of the bytes in the *7-byte-chunk* occurred **three times or more** *sequentially*. Bytes of equal value within the *7-byte-chunk* but are immediately preceded and followed by bytes of differing values, as well as like bytes that appear in pairs **do not** qualify for this case.

For example, consider the following streams:

$$(\texttt{0x04}, \texttt{0x11}, \texttt{0x11}, \texttt{0x11}, ...), \tag{1}$$

$$(\texttt{0x04}, \texttt{0x11}, \texttt{0x04}, \texttt{0x51}, ...), \tag{2}$$

and

$$(\texttt{0x6C}, \texttt{0x39}, \texttt{0x6C}, \texttt{0x6C}, \texttt{0x44}, ...). \tag{3}$$

The first stream (Eq. 1) would include the byte `0x11` with a quantity of **3** as the second byte type of its 7-byte-chunk. This stream would qualify for this case.

The second stream (Eq. 2) would include the byte `0x04` with a quantity of **1** as the first byte type of its 7-byte-chunk, and would include the (same) byte `0x04` with a quantity of **1** as the third byte type of its 7-byte-chunk. This stream would **not** qualify for this conditional case (see case 2 where it would qualify).

The third stream (Eq. 3) would include the byte `0x6C` with a quantity of **1** as the first byte type of its 7-byte-chunk, the byte `0x39` with a quantity of **1** as the second byte type, and the byte `0x6C` with a quantity of **2** as the third byte type of its 7-byte-chunk. Note how the bytes `0x6C` are **not** aggregated into one. This stream would **not** qualify for this conditional case (see case 2 where it would qualify).

This case uses a header byte to provide information about the following set of bytes. This header byte is a full header byte, meaning that all the bits of the header are used to convey information other than the data. This header byte has the following encoding:

$$\overset{\displaystyle\overbrace{\phantom{\_\ \_\ \_\ \_\ \_\ \_\ \_}}^{\substack{\text{Repeated Byte}\\ \text{Indicators}}}}{\underset{\underset{\text{Compression Type Bit}}{\nearrow}}{0}\ \underset{1}{\_}\ \underset{2}{\_}\ \underset{3}{\_}\ \underset{4}{\_}\ \underset{5}{\_}\ \underset{6}{\_}\ \underset{7}{\_}} \tag{4}$$

The **Compression Type Bit** indicates that the following set of bytes is encoded using the case-1 sub-algorithm.

Each bit of the **Repeated Byte Indicators** encode whether the byte at that positional index is a sequential byte or not.

If the bit at position 1 (on Eq. 4) is set to "0", then the byte is an individual byte and does not have any multiplicity. Otherwise, if the bit is set to "1", then this indicates that the byte *does* have multiplicity and the pair of bytes at that relative index should be considered: The first byte of the pair is the actual value of the byte. The second byte of the pair is the number of times this byte appears sequentially in the original data. This repeats for all the other indices of *repeated bit indicators*.

For example, the stream

$$(\texttt{0x03, 0x74, 0x04, 0x04, 0x35, 0x35, 0x64, 0x64, 0x64, 0x64, 0x00, 0x00, 0x00, 0x00, 0x00, 0x56}) \tag{5}$$

would become the *7-byte-chunk* represented by the array of 2-tuples where the first value of each 2-tuple represents the byte value, and the second value represents the multiplicity of that byte (the number of times it occurred in a row). The position of the 2-tuple in the array represents position of that byte relative to the other 2-tuples in the *7-byte-chunk*. As follows:

$$((\texttt{0x03}, 1), (\texttt{0x74}, 1), (\texttt{0x04}, 2), (\texttt{0x35}, 2), (\texttt{0x64}, 4), (\texttt{0x00}, 5), (\texttt{0x56}, 1)) \tag{6}$$

when encoded back into a byte array, using this algorithm, we would get:

$$(\texttt{0b00011110, 0x03, 0x74, 0x04, 0x00, 0x35, 0x00, 0x64, 0x02, 0x00, 0x03, 0x56}). \tag{7}$$

This resulting compressed data is 12 bytes in length, whereas the original stream of data was 16 bytes in length. That's a 25% decrease, and thus a successful compression.

You may notice that in the compressed stream of bytes (Eq. 7) that the multiplicity bytes do not match the values found in the corresponding 7-byte-chunk (Eq. 6). In fact, they are all a value of 2 lower than what might have been expected. This is because in the header byte, we already indicate when a pair of bytes will be present, and this only occurs when the byte value has a multiplicity of at least 2. Thus, we can set the multiplicity byte to start counting at 2 (represented by the value 0) so that we can store a byte with a multiplicity of up to 257 instead of only up to a multiplicity of 255; slightly increasing our potential for compression.

### 3.1.1 Worst Case Scenario

For this case of the algorithm to run, the *7-byte-chunk* must at least have one sequential byte with a multiplicity of 3. In the worst case scenario, there is one singular byte with a multiplicity of 3, and all other bytes have a multiplicity of 1 or 2.

The first byte outputted is the header info byte, this increases the output buffer length by 1.

For all the bytes of multiplicity 1, they are not encoded, and their byte is simply re-added to the stream at their appropriate index. Thus, for each of these bytes, the length of the output remains the same.

For all the bytes of multiplicity 2, they first output the byte of their own value, then followed by the byte encoding the quantity of bytes condensed. Since two bytes were encoded back into two bytes, the length of the output remains the same.

For the singular byte of multiplicity 3, this byte is encoded just the same as was done for bytes of multiplicity 2, except the quantity byte has a value of 3. Thus, we encoded 3 bytes from the input down to 2 bytes in the output, resulting in an output buffer length decrease of 1.

The decrease of 1 byte from the singular multiplicity 3 byte counteracts the increase of 1 byte from the header info byte, resulting in a final output buffer length equal to that of the input buffer length.

Thus, even in the worse case scenario, this component of the algorithm does not produce a compressed data buffer with a length greater than the uncompressed data buffer.

## 3.2 Case 2: 2 or Less Sequential Bytes Found

This case occurs when all the bytes in the *7-byte-chunk* are of multiplicity 1 or 2.

Since there are no sequential bytes from the 7-byte-chunk that can be compressed by encoding their multiplicity, this sub-algorithm uses the fact that the data value from every byte of the source data is only from `0x00` to `0x7F` to *condense* the data by removing the unused high-order-bit from every data byte. This sub-algorithm becomes more efficient when it can condense more bytes at once so it can aggregate all the "saved" high-order-bits together into the saving of an entire byte. To do this, after meeting the condition for using this sub-algorithm from the contents of the *7-byte-chunk*, it will look more bytes ahead to determine if there are more bytes with a multiplicity of less than three that can be encoded into this byte set with the desired result of saving an entire byte.

Like case 1, this case also uses a header byte to provide information about the following set of bytes. This header byte is a *partial* header byte, meaning that a portion of the bits are used to encode information, and the rest are used to store data.

This header byte is encoded as follows:

$$\underset{\underset{\text{Compression Type Bit}}{\nearrow}}{1}\,\underset{0\ \ 1}{\underbrace{\_\,\_}_{}}\,\underset{}{\boxed{\_\_\_\_\_}} \tag{8}$$

Extension Code     Data Bits

The **Compression Type Bit** indicates that the following set of bytes is encoded using the case-2 sub-algorithm.

The **Extension Code** is the encoding that indicates how many *extra* bytes will be included in this byte set. Since only certain quantities of input bytes accumulate to result in saving a whole byte, these encodings map to specific values (see below).

The **Data Bits** are the hight order bits of the first byte of data that is encoded of the following byte sequence.

### 3.2.1 Extension Codes

| Extension Code | | # Input | # Output | Trailing | Explanation |
|---|---|---|---|---|---|
| Bit 0 | Bit 1 | Bytes | Bytes | Spare Bits | |
| 0 | 0 | 7 | 7 | 4 | This is the default encoding. This is the encoding that insures that this algorithm does not increase the number of bytes in the resulting output buffer compared to the length of the input buffer. |
| 0 | 1 | 11 | 10 | 0 | This is the number of input bytes required to save **1** whole byte in the output. |
| 1 | 0 | 18 | 16 | 0 | This is the number of input bytes required to save **2** whole bytes in the output. |
| 1 | 1 | 25 | 22 | 0 | This is the number of input bytes required to save **3** whole bytes in the output. |

### 3.2.2 Encoding